

Process Management

(Introduction, Process, Threads & CPU Scheduling)

Process Management-1

Address translation

- (a) DMA for disk transfer
- (b) least two modes of CPU execution
- (c) privileged and non-privileged
- (d) Demand paging

[1999 : 2 M]

- 1.1** Which scheduling policy is most suitable for a time-shared operating systems?
- Shortest Job First
 - Round Robin
 - First come first serve
 - Elevator

- 1.2** The sequence is an optimal non-preemptive scheduling sequence for the following jobs which leaves the CPU idle for unit(s) of time.

- 1.5** Four jobs to be executed on a single processor system arrive at time 0⁺ in the order A, B, C, D. Their burst CPU time requirements are 4, 1, 8, 1 time units respectively. The completion times are 4, 1, 8, of one time unit is
- 10
 - 4
 - 8
 - 9

[1996 : 2 M]

- 1.6** Which of the following is an example of a spooled device?
- The terminal used to enter the input data for the C program being executed
 - An output device used to print the output of a number of jobs.
 - The secondary memory device in a virtual storage system
 - The swapping area on a disk used by the swapper.

[1998 : 1 M]

- 1.7** Consider n processes sharing the CPU in a round-robin fashion. Assuming that each process switch takes s seconds, what must be the quantum size q such that the overhead resulting from process switching is minimized but at the same time each process is guaranteed to get its turn at the CPU at least every t seconds?

$$\begin{array}{ll} \text{(a)} & q \leq \frac{t - ns}{n - 1} \\ \text{(b)} & q \geq \frac{t - ns}{n - 1} \\ \text{(c)} & q \leq \frac{t - ns}{n + 1} \\ \text{(d)} & q \geq \frac{t - ns}{n + 1} \end{array}$$

[1998 : 1 M]

- 1.8** Consider a set of n tasks with known runtimes r_1, r_2, \dots, r_n to be run on a uniprocessor machine. Which of the following processor scheduling algorithms will result in the maximum throughput?

[1998 : 1 M]

- 1.9** System calls are usually invoked by using
- a software interrupt
 - polling
 - an indirect jump
 - a privileged instruction

[1999 : 1 M]

- 1.10** A multi-user, multi-processing operating system cannot be implemented on hardware that does not support

A device

- (b) Timer
- (c) Scheduler process
- (d) Power failure

[2001 : 2 M]

- 1.7** Which combination of the following features will suffice to characterize an OS as a multi-programmed OS? (i) more than one program may be loaded into main memory at the same time for execution. (ii) If a program waits for certain events such as I/O, another program is immediately scheduled for execution. (iii) If the execution of a program terminates, another program is immediately scheduled for execution
- (i)
 - (ii) and (iii)
 - (i) and (ii)
 - (i) and (iii)

[2002 : 2 M]

- 1.11** Which of the following actions is/are typically not performed by the operating system when switching context from process A to process B?
- Saving current register values and restoring saved register values for process B.
 - Changing address translation tables.
 - Swapping out the memory image of process C.
 - To the disk.

[1999 : 2 M]

- 1.12** A processor needs software interrupt to
- Test the interrupt system of the processor
 - Implement co-routines
 - Obtain system services which need execution of privileged instructions
 - Return from subroutine

[2001 : 1 M]

- 1.13** A CPU has two modes: privileged and non-privileged. In order to change the mode from privileged to non-privileged
- a hardware interrupt is needed
 - a software interrupt is needed
 - a privileged instruction (which does not generate an interrupt) is needed
 - a non-privileged instruction (which does not generate an interrupt) is needed

[2001 : 1 M]

- 1.14** Consider a set of n tasks with known runtimes r_1, r_2, \dots, r_n to be run on a uniprocessor machine. Which of the following processor scheduling algorithms will result in the maximum throughput?

[2003 : 2 M]

- 1.15** Consider the following scheduling algorithms is non-preemptive
- Round-Robin
 - Shortest-Job-First
 - Highest-Response-Ratio-Next
 - First-Come-First-Served

[2001 : 1 M]

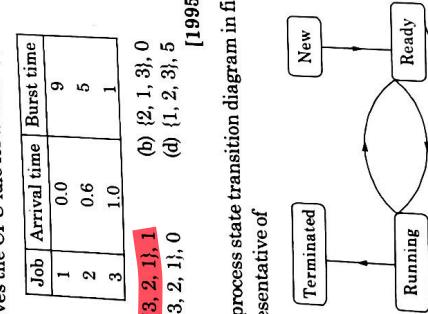
- 1.16** Which of the following statements are true?
- User-level threads are transparent to the Kernel
 - Context switch is faster with Kernel-supported threads
 - For user-level threads, a system call can block the entire process
 - Kernel-supported threads can be scheduled independently

[2004 : 1 M]

- 1.17** Which combination of the following features will suffice to characterize an OS as a multi-programmed OS? (i) more than one program may be loaded into main memory at the same time for execution. (ii) If a program waits for certain events such as I/O, another program is immediately scheduled for execution. (iii) If the execution of a program terminates, another program is immediately scheduled for execution
- (i)
 - (ii) and (iii)
 - (i) and (ii)
 - (i) and (iii)

[2002 : 2 M]

- 1.18** Draw the process state transition diagram of an OS in which (i) each process is in one of the five states: created, ready, running, blocked (i.e. sleep or wait), or terminated, and (ii) only non-preemptive scheduling is used by the OS. Label [2002 : 2 M]



- 1.19** Consider a set of n tasks with known runtimes r_1, r_2, \dots, r_n to be run on a uniprocessor machine. Which of the following processor scheduling algorithms is faster with respect to user-level threads and kernel-supported threads

[2003 : 2 M]

- 1.20** Consider the following statements with respect to user-level threads and kernel-supported threads
- Context switch is faster with Kernel-supported threads
 - For user-level threads, a system call can block the entire process
 - Kernel-supported threads can be scheduled independently
 - User-level threads are transparent to the Kernel

[2004 : 1 M]

(b) Ready to Waiting

(a) Waiting to Running

is are NOT possible?

Waiting

to

Runnning

Which of the following processes is/are possible

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

?

All the threads share address space but other entities like, stack, PC, registers are not shared and every thread will have its own.

Generally every thread of a process have their variables are shared by every thread of a process. Both heap and global own PC and stack. Both heap and global variables are shared by every thread of a process.

Using preemptive SRTF algorithm Gantt chart will be,

Process	Arrival time	Burst time	Completion time	Average turn around time = $\frac{33}{4} = 8.25$	
				Number of child processes created	Average turn around time = $\frac{33}{4} = 8.25$
1	0	10	20	1	1
2	3	7	20	2	2
3	6	10	30	3	3
4	10	13	43	4	4
5	13	5	48	5	5

To minimize the average waiting time, we need to select the shortest remaining time process, first, because all are arriving at the same time, and they have unequal CPU burst times. All other options will not minimize the waiting time. So, the answer is SRTF algorithm.

To minimize the average waiting time, we need to select the shortest remaining time process, first, because all are arriving at the same time, and they have unequal CPU burst times.

Average turn around time is 8.25.

Assume, $Z = 2$

P ₁	P ₂	P ₃	P ₁	P ₃	Waiting time	Completion time	Arrival time
0	1	2	3	4	6	9	0
1	2	3	4	5	6	10	1
2	3	4	5	6	7	12	2
3	4	5	6	7	8	14	3
4	5	6	7	8	9	16	4
5	6	7	8	9	10	18	5
6	7	8	9	10	11	20	6

WT	CT	P.No.	AT	BT
0	10	A	0	16
0	26	B	0	20
26	46	C	0	10
10	26		0	10

WT	CT	P.No.	AT	BT
0	10	A	0	16
0	26	B	0	20
26	46	C	0	10
10	26		0	10

WT	CT	P.No.	AT	BT
0	10	A	0	16
0	26	B	0	20
26	46	C	0	10
10	26		0	10

WT	CT	P.No.	AT	BT
0	10	A	0	16
0	26	B	0	20
26	46	C	0	10
10	26		0	10

WT	CT	P.No.	AT	BT
0	10	A	0	16
0	26	B	0	20
26	46	C	0	10
10	26		0	10

WT	CT	P.No.	AT	BT
0	10	A	0	16
0	26	B	0	20
26	46	C	0	10
10	26		0	10

WT	CT	P.No.	AT	BT
0	10	A	0	16
0	26	B	0	20
26	46	C	0	10
10	26		0	10

WT	CT	P.No.	AT	BT
0	10	A	0	16
0	26	B	0	20
26	46	C	0	10
10	26		0	10

WT	CT	P.No.	AT	BT
0	10	A	0	16
0	26	B	0	20
26	46	C	0	10
10	26		0	10

WT	CT	P.No.	AT	BT
0	10	A	0	16
0	26	B	0	20
26	46	C	0	10
10	26		0	10

WT	CT	P.No.	AT	BT
0	10	A	0	16
0	26	B	0	20
26	46	C	0	10
10	26		0	10

WT	CT	P.No.	AT	BT
0	10	A	0	16
0	26	B	0	20
26	46	C	0	10
10	26		0	10

WT	CT	P.No.	AT	BT
0	10	A	0	16
0	26	B	0	20
26	46	C	0	10
10	26		0	10

WT	CT	P.No.	AT	BT
0	10	A	0	16
0	26	B	0	20
26	46	C	0	10
10	26		0	10

WT	CT	P.No.	AT	BT
0	10	A	0	16
0	26	B	0	20
26	46	C	0	10
10	26		0	10

WT	CT	P.No.	AT	BT
0	10	A	0	16
0	26	B	0	20
26	46	C	0	10
10	26		0	10

WT	CT	P.No.	AT	BT
0	10	A	0	16
0	26	B	0	20
26	46	C	0	10
10	26		0	10

WT	CT	P.No.	AT	BT
0	10	A	0	16
0	26	B	0	20
26	46	C	0	10
10	26		0	10

WT	CT	P.No.	AT	BT
0	10	A	0	16
0	26	B	0	20
26	46	C	0	10
10	26		0	10

WT	CT	P.No.	AT	BT
0	10	A	0	16
0	26	B	0	20
26	46	C	0	10
10	26		0	10

WT	CT	P.No.	AT	BT
0	10	A	0	16
0	26	B	0	20
26	46	C	0	10
10	26		0	10

WT	CT	P.No.	AT	BT

<tbl_r cells="5" ix="4

- (a) Both I and II are true.
 (b) I is true but II is false.
 (c) II is true but I is false.
 (d) Both I and II are false.

2.15 Two concurrent processes P1 and P2 use four shared resources R1, R2, R3 and R4, as shown below:

Computer: P1: P2:
 Use R1; Use R2;
 Use R2; Use R3;
 Use R3; Use R4;

2.16 The code does not implement a binary semaphore correctly. Which one of the following is true?

- (a) Both I and II are false.
 (b) I is true but II is false.
 (c) II is true but I is false.
 (d) Both I and II are true.

2.17 The above implementation of barrier is incorrect

2.18 Which one of the following rectifies the problem in the implementation?

- (a) Lines 6 to 10 are simply replaced by

the implementation?

2.19 Following routine to achieve mutual exclusion for readers-writers problem, in the following incomplete code for semaphores, the following互換化 (mutual exclusion) can be achieved through multiple readers and writer are used to obtain synchronization

2.20 Following routine to achieve mutual exclusion for processes P1 and P2 use critical flag in the set arrives at the barrier and waits for processes in the set to leave the barrier and S is released. Consider the following C implementation of V is wrong.

2.21 Synchronization in the classic readers and writers problem can be achieved through multiple readers-writers problem, in the following incomplete code for semaphores, the following互換化 (mutual exclusion) can be achieved through multiple readers and writer are used to obtain synchronization

2.22 The P and V operations on counting semaphores, where s is a counting semaphore, are defined as

2.23 Two processes, P1 and P2, need to access a critical section of code. Consider the following

2.24 The variables process-arrived and process-left are

2.25 Consider the following statements.

2.26 If readcount = 0 then S4

readcount = readcount - 1

reading is performed

2.27 wait (wrt)

signal (wrt)

wait (wrt)

signal (wrt)

wait (wrt)

signal (wrt)

wait (wrt)

part (v);

Release L;

$y = y + 1;$

$x = x + 1;$

Acquire L;

$int y = 0;$

print(x);

foo();

main();

Lock L;

// global

int x = 0;

processes

by two

processes

P₁ and P₂

and each

of the

segments

(in a mix of C and pseudo-code), invoked

by the

same

shared

variable X.

The initial value of X is 10.

Suppose each invocation decrements one and

shares

the same

variable

semaphores

initially

set to 1.

create a thread to execute foo(); // Thread T₁

create a thread to execute foo(); // Thread T₂

wait for the two threads to finish execution;

int total = 0;

for (int i = 0; i < 10; i++)

total += i;

print(total);

return total;

}

else

if (P₁ is owned by another process P then

Acquire R

lock L

Acquire Lock L // a global lock

unlock OwnerResource (Resource R)

instance: The following scheme is used to own a resource

resources are done by holding a global lock (L).

processes at a time. Owing and freeing of

resources can be owned by only one

type. Each instance per resource

resource types, with one instance per resource

and multiple shared

resources are held by the same

process.

correct;

consider the following code snippet using the fork() and wait() system calls. Assume that the code complies

with the following statement:

At least one of P₁ and P₂ will print the value of x as 4,

both T₁ and T₂ will print the value of x as 1,

fork();

print("Hello");

wait("Hello");

X--;

print("World");

fork();

print("Hello");

wait("Hello");

print("World");

X--;

print("Hello");

fork();

print("Hello");

wait("Hello");

X--;

print("World");

X--;

print("Hello");</

